

Edge / Local (Air-gapped)

M1 (MT-LNN) Engine

Personal State Capsule (4.1KB h_{prev})
Lifelong memory log

Predictive Error Monitor (Active UI/UX)

Inference Loop

10s Silent

Deep Decoder

(For Queries)

Deep Encoder

(For Responses)

CPU / NPU inference only

M1 (MT-LNN) Engine

- Personal State Capsule (4.1KB h_prev)\nLifelong memory
- Predictive Error Monitor (Active UI/UX)
- CPU / NPU inference only

State Capsule (4.1KB h_prev)\nLifelong men

Predictive Error Monitor (Active UI/UX)

10s Silent Circle

1. Submits

CPU / NPU inference only

```
graph TD; A[Vector DB FactsRAG] --> B[Gemini 3.1 / GPT-4o LLM API];
```

Cloud (Knowledge Oracle)

Vector DB (Facts/RAG)

Gemini 3.1 / GPT-4o (LLM API)

Summa raw facts

Vector DB (Facts/RAG)

emini 3.1 / GPT-4o\n(LLM AP

The diagram illustrates the Gemini API workflow. It starts with a user input box containing the text "What is the capital of France?". This input is sent to the "Gemini API" box, which then sends a "Query" to the "Gemini" box. The "Gemini" box returns a "Response" to the "Gemini API" box, which then displays the "Response" in the "Output" box. The "Output" box contains the text "Paris".