

MT-LNN: Microtubule-Inspired Liquid Neural Network

Shattering the Memory Wall for the Next Generation of Long-Context AI

Investor Pitch Deck

Awareness

May 2026

Paper: <https://huggingface.co/EverestAn/MT-LNN>

GitHub: <https://github.com/everest-an/M1>



- ▶ **1. Symbolism:** Logical and visual reasoning, closest to human conscious thought.
- ▶ **2. Bayesianism:** Prior/posterior probability reasoning via Bayesian networks.
- ▶ **3. Analogizer:** SVMs and nearest-neighbor classification.
- ▶ **4. Evolutionism:** Mutational and evolutionary algorithms.
- ▶ **5. Connectionism (Current Mainstream):** Neural networks, Large Models, XAI, and **Liquid Neural Networks (LNN)**.
- ▶ **6. Reinforcement Learning:** Reward-punishment optimization.
- ▶ **7. Multi-Agent Systems:** Simulating interaction between multiple intelligent agents.
- ▶ **8. Continual Learning:** AI evolving continuously without total retraining.



- ▶ **The Problem:** Large Language Models (LLMs) are hitting a physical “Memory Wall”. Scaling past 200K tokens leads to exponential costs and severe Out-Of-Memory (OOM) constraints during deployment.
- ▶ **The Solution:** MT-LNN replaces static Transformer mechanisms with a continuous-time Liquid Neural Network inspired by human brain microtubules.
- ▶ **The Edge:** Constant $O(1)$ memory footprint, eliminating the KV Cache dependency. High performance in ultra-long context reasoning ($> 128K$ tokens) on consumer-grade hardware.
- ▶ **The Ask:** Seeking strategic funding to scale our models from the 1.5B component level to a fully native 7B Enterprise architecture.

The AI Industry's Memory Crisis



- ▶ **The Cost of Long Context:** The AI industry is in an arms race for 200K+ token windows, but the underlying computational and memory costs are exploding.
- ▶ **OOM (Out-of-Memory) Limits:** For on-premise and local deployments, the true bottleneck isn't logic or processing power—it is strictly physical VRAM limits.
- ▶ **The Compute Trap:** Cloud providers are forced to stack phenomenally expensive GPU clusters to handle concurrent, long-context user requests due to legacy architecture flaws.

The Root Cause: KV Cache Bottleneck




- ▶ **A Structural Limitation:** The standard Transformer architecture operates like an exhaustive recorder; predicting the next token requires holding all previous tokens' features in memory (KV Cache).
- ▶ **Complexity Disaster:** As context grows, memory complexity scales linearly $O(N)$, but compute complexity scales quadratically $O(N^2)$.
- ▶ **Physical Constraints:** Processing a 100,000-word document instantly fills an expensive A100 GPU. Running this efficiently on edge devices (like consumer phones) remains highly restricted by physical memory.

Cost Explosion vs. Perfect Scaling




- ▶ Transformer family models depend on growing KV cache, leading to rising serving cost and memory pressure under long-context traffic.
- ▶ MT-LNN uses a compact recurrent latent state and microtubule-inspired selective dynamics to control memory growth.
- ▶ Product implication: lower deployment friction for private/on-prem and edge scenarios.



Biological Foundation: What are Microtubules?



- ▶ **Sub-cellular Processing:** While conventional AI mimics the network of neurons, actual biological cognition is regulated at a deeper level within the **Microtubules**—the cell's cytoskeleton.
- ▶ **Quantum-Level Dynamics:** Microtubules process information through continuous, dynamic protein states rather than binary synapses, enabling higher-dimensional fluid data integration.
- ▶ **The Pathway to LNN:** This biological flow directly inspires the **Liquid Neural Network (LNN)**, creating an architecture that inherently adapts to time-series data without rigid matrices.





- ▶ **Refusing to Store Every Pixel:** The human brain does not—and physically cannot—memorize the exact pixels of every scene spanning a decade.
- ▶ **Working Memory & Selective Forgetting:** Human cognition relies heavily on “Working Memory” (retaining core logic) and “Selective Forgetting” (discarding irrelevant noise).
- ▶ **AI’s Blind Spot:** Current LLMs are purely vast matrix multiplications lacking this dynamic time-flow mechanism or a Global Workspace Theory (GWT) path for conscious-level integration.


What is a Liquid Neural Network?



- ▶ **Dynamic Flowing States:** Liquid Neural Networks (LNNs) maintain an internal state that fluidly adapts to new data over time, contrasting sharply with static Transformer weights.
- ▶ **Compression and Abstraction:** The core mechanism: extract the essential synopsis, forget redundant filler, and compress thousands of words into a high-dimensional, time-evolving hidden state.
- ▶ **Brain-Like Foundation:** It is the first architecture to physically and dynamically mimic how microtubule structures process information in biological cells.



- ▶ **The Probabilistic Illusion:** Traditional Transformers can perfectly predict the sequence “*the **apple** fell from the tree to the **ground**.*” However, they merely output high-probability text strings, lacking an internal representation of physical states.
- ▶ **Causal Physical Perception:** MT-LNN breaks this limitation. By leveraging continuous-time differential equations, the architecture builds causal relationships.
- ▶ **Simulating Reality:** It dynamically evolves its internal hidden states, fundamentally grasping the physical process of the **apple falling** over time rather than just memorizing word frequencies.




MT-LNN Dynamic Concept Positioning

The MT-LNN Architecture



Microtubule Liquid Neural Network:

- ▶ **Replacing the Cache Wall:** We completely discard the traditional $O(T)$ KV Cache, introducing an $O(1)$ **Working Memory Decay** matrix. Context size scales infinitely with zero added memory footprint.
- ▶ **Replacing Static FFNs:** The architecture integrates 13 parallel, continuous-time liquid network pathways equipped with *endogenous compute skipping*.
- ▶ **Native Conscious Computing:** The world's first AI architecture to break the "Anesthesia Benchmark," capable of generating language while tracking an inherent cognitive collapse metric ($\hat{\Phi}$).





- ▶ **Predictive Coding System:** Moving beyond pure “next word guessing”, MT-LNN introduces inter-scale predictive coding. Abstract reasoning channels natively self-supervise faster perception channels, forcing deep logic comprehension.
- ▶ **Endogenous Compute Skipping:** Utilizing dynamic κ logic gates, the architecture autonomously shuts off dormant channels mid-computation, yielding **exponential ROI cost-savings** on inference chips.
- ▶ **Explainable Causal Logging:** Say goodbye to black boxes. MT-LNN state matrices are natively parsed by *Causal Chain* and *Self-Monitor* extraction heads, rendering internal logical stepping completely observable.



- ▶ **Hyper-Specialization:** We dedicate architecture to complex reasoning rather than storing vast arrays of encyclopedia trivia.
- ▶ **Laser-Focused Optimization:** Parameter capacity is heavily invested in *hardcore logical reasoning* and *massive cross-context structural memory constraints*.
- ▶ **Enterprise Application:** Perfect for delivering world-class semantic extraction on a 50,000-word financial audit or a massive legacy codebase branch, strictly operating on local secure hardware.


Core Performance 1: Selective Copy & Noise Resistance



Surviving Massive Noise ($T = 229$):

- ▶ **Mainstream Collapse:** When timeline tracking expands and noise floods the context, traditional model retention plummets to a mere 2% random-guess baseline.
- ▶ **MT-LNN High Robustness:** By autonomously filtering what to retain and what to prune via liquid states, our precise match rate remains stable at **96.5%**.

At an ultra-compact scale ($\sim 200K$ parameters), MT-LNN outperforms mainstream architectures by a +94% accuracy margin under heavy noise.



Core Performance 2: 128K Needle in a Haystack



- ▶ **Solving the Focal Illusion Crisis:** We eradicate the "Lost in the Middle" syndrome that routinely cripples the industry. Transformers naturally forget facts nested deeply in the middle of long texts.
- ▶ **Penetrating the 128,000-Word Ocean:** Drop a microscopic piece of data into a boundless sea of tokens. Whether parked at the beginning, the deepest core, or the very end, MT-LNN retrieves it with a linear **99%+ accuracy**.


The Ultimate Vision: Awareness Network

(Hybrid Edge-State Intelligence System)



- ▶ **Strategic Disruption:** Rather than relying on legacy Transformer APIs like Gemini and GPT-4o, we are vertically integrating our breakthrough MT-LNN architecture into an end-to-end ecosystem.
- ▶ **Core Philosophy:** The Edge **AwareLiquid Engine** serves as the continuous "soul" and reasoning engine, while our proprietary **Awareness Cloud** (High-Param MT-LNN) acts as the global knowledge Oracle.
- ▶ **Hybrid Edge-State Architecture:** The Awareness Network runs purely local logic natively, routing to the Awareness Cloud only for factual retrieval while maintaining absolute data sovereignty.
- ▶ **Privacy by Physics:** Core cognitive state remains perpetually contained within a local 4.1 KB $O(1)$ State Capsule, completely severing the "harvest and charge" cycle of legacy providers.

Awareness Network Architecture



The End of the "Token-Burning" Compute Arms Race



▶ The Legacy "Compute Trap":

- ▶ The Transformer paradigm forces every interaction to re-transmit and re-compute massive context histories, locking the industry into an endless "Token-Burning" death spiral.

▶ Our Paradigm: Brain-like Edge Engine + Cloud Knowledge Base

- ▶ **Local Brain (AwareLiquid):** 90% of logical reasoning, continuous memory, and private deliberation are handled purely on-device via a lightweight MT-LNN, achieving **Zero-Cost Local Compute**.
- ▶ **Awareness Cloud:** The cloud is no longer a high-frequency token black hole. It serves simply as a heavy-parameter knowledge base.

▶ The Business Reality: Escaping the Compute Wars

- ▶ Local compute resolves the vast majority of tasks. Ultra-lightweight knowledge retrieval is triggered *only* for factual blind spots—making scalable AI finally free from server-farm burn rates.

Four Pillars of the Experiential Shift



- ▶ **1. Latent Reasoning Loop (Silent Deep Deduction):**
 - ▶ **Cloud AI:** Fakes reasoning by spitting out verbose Chain-of-Thought (CoT) tokens.
 - ▶ **M1:** Embarks on 10 seconds of "silent calculation" running internal ODEs natively. It strikes truth without bloated text, exhibiting wise, human-like intuition.
- ▶ **2. Lifelong Personal State Capsule:**
 - ▶ **Cloud AI:** Amnesic per session, requiring external Vector RAG.
 - ▶ **M1:** Compresses millions of inputs into a constant h_{prev} . Load the capsule onto a new device, and it instantly inherits years of synchronized synergy.
- ▶ **3. Predictive Error Monitor (Active AI):**
 - ▶ Shifts from passive Q&A to active prediction. It anticipates intent top-down and triggers red-light interventions the moment user logic violates globally established context.
- ▶ **4. Cloud Oracle Router:**
 - ▶ In factual blind spots, AwareLiquid dispatches surgical queries to our proprietary **Awareness Cloud**. Legacy APIs are excluded; instead, AwareLiquid fluidly synthesizes world knowledge through its personalized state via a unified MT-LNN ecosystem.



- ▶ **1. Extreme Cost Compression ($O(1)$ Memory):**
 - ▶ **Transformer:** 1000-token KV Cache bloats to 1020 KB. Memory costs scale boundlessly.
 - ▶ **M1:** The Personal State Capsule strictly bounds h_{prev} memory to **4.1 KB**, executing zero-cost inference via Edge CPU/NPU processors.
- ▶ **2. True Accuracy (Dynamic "Needle in a Haystack"):**
 - ▶ **Transformer:** Selective Copy recovery decays rapidly to **2.3%**.
 - ▶ **M1:** Sustains **96.5%** accuracy (42x better). Biological Predictive Coding inherently suppresses Hallucinations via local consistency checks.
- ▶ **3. Speed (Zero-Latency & Resonance):**
 - ▶ **Zero Network Overhead:** Bypasses transmitting 100k+ history tokens to cloud API servers.
 - ▶ **Sparse Resonance Acceleration:** Activating microtubule scale-gates yields an immediate **13% CPU processing speedup** (from 6528 to 7390 Tok/s).



- ▶ **1. Memory Compression Evidence (1000-token footprint)**
 - ▶ **Data:** Transformer KV = 1020 KB v.s. MT-LNN = 4.1 KB.
 - ▶ **Meaning:** Enables continuous "always-on" AI on strictly limited edge devices without RAM bottlenecks (OOM).
 - ▶ **Reproduce Action:** Run `python benchmarks/operator_compression_report.py`.
- ▶ **2. Long-Context Accuracy Evidence ("Needle in a Haystack")**
 - ▶ **Data:** In 200K parameter head-to-head, Transformer retrieval crashes to 2.3% while MT-LNN sustains 96.5% (a 42x advantage).
 - ▶ **Meaning:** Absolute fidelity for tracking user history, inherently suppressing hallucinations.
 - ▶ **Reproduce Action:** Run `python benchmarks/run_benchmark.py`.
- ▶ **3. CPU Speed Evidence (Sparse Resonance)**
 - ▶ **Data:** Dense processing at 6528 Tok/s jumps to 7390 Tok/s.
 - ▶ **Meaning:** Biologically-inspired "sleep" patterns skip redundant compute, slashing inference bills.
 - ▶ **Reproduce Action:** Run `python benchmarks/sparse_resonance_ablation.py`.

Market Sizing (TAM/SAM/SOM)



- ▶ **TAM (Total Addressable Market) - \$120B+:** The global generative AI and LLM inference market spanning cloud services, enterprise applications, and edge intelligence.
- ▶ **SAM (Serviceable Addressable Market) - \$45B:** Enterprise data processing, legal review, financial audits, and local-first AI software markets where data privacy is paramount.
- ▶ **SOM (Serviceable Obtainable Market) - \$2.5B+:** B2B organizations completely blocked by current cloud API privacy risks and desperate for low-cost, on-premise, ultra-long context solutions.



- ▶ **1. B2B Enterprise Licensing (On-Premise):**
 - ▶ Selling direct private licenses to heavily-regulated industries (Law, Finance, Defense).
 - ▶ Deployment straight to their internal servers. 100% localized, 0% data leak risk.
- ▶ **2. High-Efficiency Cloud API:**
 - ▶ Because our inference costs are $\sim 90\%$ lower than Transformer APIs, we can undercut market leaders significantly while preserving massive profit margins.





- ▶ **Phase 1: Open-Source Community Anchor:** Release the 1.5B components to OSS. Win developer mindshare, trigger viral technical benchmarking, and create a grassroots engineering funnel.
- ▶ **Phase 2: Targeted B2B Pilot Deployments:** Launch private proofs-of-concept (PoCs) with 5 Tier-1 law firms and venture funds. Prove ROI purely on the hours saved from document parsing.
- ▶ **Phase 3: The API Marketplace:** Roll out a self-serve developer portal for high-volume, cheap, long-context inference wrappers substituting standard OpenAI/Anthropic pipelines.

Competitive Landscape Matrix



Features	MT-LNN	GPT-4/Claude	Mamba/RWKV	Local Llama
Ultra-Long Context	Yes	Yes	Degrades	Fails (OOM)
Privacy / On-Prem	Yes	No (Cloud)	Yes	Yes
Low Hardware Req.	Yes (MacBook)	No (A100)	Yes	No
Constant Memory $O(1)$	Yes	No	Yes	No
Conscious Core Eval	Yes	No	No	No

Market Positioning & Concurrency Competitiveness





- ▶ **Solving the Hardware Bottleneck:** We are migrating advanced models directly into highly-constrained hardware: wearable AR glasses, native automotive local hubs, and flagship smartphones.
- ▶ **True “Always-On” Intelligence:** Our architecture supports continuous real-time audio/video processing. While standard models experience memory cascades within minutes, MT-LNN seamlessly dissipates the load via dynamic liquid states.



- ▶ **The Core Client Persona:** Top-tier law firms reviewing 1,000-page case files; hedge funds auditing 10,000 corporate annual reports; code reviewers untangling decades-old monolithic software structures.
- ▶ **100% Pain-Point Overlap:** They suffer a trifecta of issues: (1) Need for massive document context, (2) Non-negotiable private local network deployments, (3) Unwillingness to buy \$200,000 GPU racks. MT-LNN dominates this exact intersection.



- ▶ **Not Just a Visionary Paper:** We have fully escaped the academic whiteboard. At this very moment, we possess a mature, running Python/Gradio UI integrating a natively functioning AI sandbox.
- ▶ **Ready to Deploy:** The codebase is fully containerized. It runs entirely offline on a standard PC or fits natively into an enterprise intranet. It is not an idea—it is a functional intelligence asset today.

Product Roadmap (18-Month Vision)



1. **Phase 1: Validation (Current):** Utilizing small scale (1.5B components), we have achieved top-tier performance on long-context benchmarks and established an initial developer community.
2. **Phase 2: Enterprise Deployment (Months 1-6):** Secure funding to train a native 7B enterprise model explicitly fine-tuned for dense RAG and local inference. Establish initial recurring B2B sales pipelines.
3. **Phase 3: The A.G.I Checkmate (Months 6-18):** Scale to an H100 orbital cluster. Introduce multimodal processing and challenge the legacy Transformer architecture on a 100B+ parameter scale across global LLM arenas.



- ▶ **Everest An – Founder & Core Architect** (AI Memory Expert)
- ▶ **Decentralized Infra & Cryptography:** Former member of IPFS Athna, leading enterprise database construction deployment.
- ▶ **Top-Tier Investment & Ecosystem:** Former Investment Manager at Outliers Fund. Led investments in L1 Chains (Dfinity, Harmony), Infra (Analog), and Smart Contract Audits (Quantstamp).
- ▶ **Global Engineering Prowess:** Multiple-time Champion at Global Ethereum Hackathons.

Financial Projections (3-5 Years)



- ▶ **Year 1:** Target \$1.5M ARR focused purely on early-adopter Enterprise Licenses (Law & Finance PoCs) and localized consulting.
- ▶ **Year 2:** Scale to \$8M+ ARR launching our disruptive high-efficiency Cloud API for mid-market B2B applications. Breaking even on compute scale.
- ▶ **Year 4-5:** \$50M+ ARR mapping. Domination of the localized and "Always-On Edge" chip integration market; vast margin retention due to 90% cheaper algorithmic operations.



Target Raise: \$3.5M

- ▶ Fast-track 7B native foundational model training run.
- ▶ Accelerate sales/go-to-market teams to capture B2B pipeline.
- ▶ **Target Milestone:** Achieve 10 Enterprise PoCs and reach \$1.5M ARR within 12 months.

Use of Funds:

- ▶ **50% Compute & R&D:** Secure H100 clusters for full-scale pre-training.
- ▶ **30% Engineering:** Hiring top systems & CUDA optimization talent.
- ▶ **20% GTM & Sales:** Direct B2B enterprise penetration.



- ▶ **A New Scaling Path:** The traditional “more compute, more scale” Transformer paradigm faces diminishing marginal returns and severe memory constraints.
- ▶ **The Dynamic Liquid Advantage:** By transitioning to a fluid, continuous-time neural architecture, we bypass the “Memory Wall” holding back extensive AI enterprise integration.
- ▶ **Join the Next Wave:** You are backing a rigorously engineered system. We have the data, the verifiable codebase, and the highly efficient cost-basis necessary to reshape edge AI and enterprise deployments.
- ▶ *Invest in Awareness. Invest in the Liquid Computing Paradigm Shift.*



Awareness

Everest An – Founder & Core Architect

Email: everest9812@gmail.com

WeChat / Telegram: EverestAn

Paper: <https://huggingface.co/EverestAn/MT-LNN>

GitHub: <https://github.com/everest-an/M1>

MT-LNN Demo: <https://huggingface.co/spaces/EverestAn/Awareness01>

Let's crush the AI Memory Wall together.